

GREATER CHENNAI TRAFFIC POLICE
VEPERY, CHENNAI – 600 007

INTERNSHIP REPORT (2023-2024)

On

TrafficGPT -

Submitted to

Tr. R. SUDHAKAR IPS
ADDITIONAL COP, GCTP

Submitted By

Krithika V (2020PECAI108)
B.Tech AI & Data Science (2020-2024)
Department of Artificial Intelligence and Data Science
Panimalar Engineering College (Autonomous)
Chennai – 600 123

Guided By

Prathibaa K & T. Sathish Babu
Traffic Experts, GCTP

ACKNOWLEDGEMENT

My sincere thanks to “**Tr. R. Sudhakar**” IPS, Additional Commissioner of Police, Greater Chennai Traffic Police. Our heartfelt gratitude for providing us with the opportunity for internship. Your leadership and commitment to excellence have been inspiring. My sincere thanks to “**Tr. K. Baskaran**”, Assistant Commissioner of Police, Greater Chennai Traffic Police. I extend my gratitude for generously sharing crucial data that significantly contributed for Our Internship.

My special thanks to “**K. Prathibaa**”, Traffic Engineering Professional, GCTP for her guidance and moral support throughout the project.

My special thanks to “**S.Gunalan**” Sub Inspector Traffic planning GCTP, “**G. Rajmohan**”, “**T. Sathish Babu**”, Technical Experts GCTP for contributing significantly to the success of my project.

My sincere thanks to “**Dr. K. MANI, M.E., Ph.D.**”, Principal, Panimalar Engineering College, for enabling me to do my Internship with great interest.

My sincere thanks to “**Dr. S. MALATHI, M.E., Ph.D.**”, Head of department, Department of Artificial Intelligence & Data Science, Panimalar Engineering College, for enabling me to do my Internship with great interest.

Sincerely,

Krithika V (2020PECAI108)

ABSTRACT

This project introduces a novel chat application designed to offer comprehensive guidance on traffic rules, laws, and updates in a conversational manner akin to human interaction. Leveraging a vector database named Chroma, textual documents related to traffic regulations are stored as vector embeddings, forming the foundation for the chatbot's understanding of context and information retrieval. At the core of the system lies Mistral-7b-instruct, a meticulously tuned and quantized language model that generates responses based on contextual cues provided by the vector embeddings. To bridge the gap between vector embeddings and natural language, sentence transformers are employed to convert embeddings into readable text, enabling seamless communication between users and the system. Employing a retrieval-augmentation-generation (RAG) process, the converted text undergoes augmentation and subsequent generation of responses, ensuring accuracy, contextual relevance, and informativeness of the chatbot's responses. Moreover, the system is built on the FastAPI framework, ensuring efficiency, scalability, and smooth integration with various platforms and applications. By combining advanced technologies and robust software engineering practices, this project aims to revolutionize the accessibility and efficacy of traffic guidance systems, ultimately contributing to safer and more informed commuting experiences.

Keywords: Traffic Guidance, Vector Database, Language Model, FastAPI, Retrieval-Augmentation-Generation (RAG), Natural Language Understanding.

CONTENTS

S.NO	TOPIC	PAGE No
1	INTRODUCTION	1
	1.1 Problem Statement	2
	1.2 Introduction	2
2	LITERATURE SURVEY	4
	2.1 Literature Survey	5
3	SYSTEM ANALYSIS	11
	3.1 Proposed Work	12
	3.2 System Analysis	16
	3.3 Existing System	18
	3.3.1 Rule-based Chatbot	18
	3.3.2 Leveraging Large Language Models for Chatbot	18
	3.4 Social Feasibility	19
	3.5 Hardware and Software Requirements	20
	3.5.1 Hardware Requirements	20
	3.5.2 Software Requirements	20

4	System Design	22
	4.1 Flow Diagram	23
5	System Architecture	24
	5.1 System Architecture	25
	5.2 Modules	27
	5.2.1 Module 1: Data Loading and Embedding conversion	27
	5.2.2 Module 2: Querying from Loaded Data	29
6	System Implementation	32
	6.1 System Implementation Requirements	33
	6.2 Python Implementation of TrafficGPT (Backend code)	36

	6.3 Knowledge Base Update Process	41
	6.4 Output of the Chat Application with a prompt and Response	43
7	Conclusions and Future Enhancements	45
	7.1 Conclusion	46
	7.2 Future Enhancements	46
8	References	48
	8.1 References	49

CHAPTER - I

INTRODUCTION

1.1 PROBLEM STATEMENT:

Despite advancements in technology, accessing comprehensive and up-to-date information on traffic regulations, laws, and updates remains a significant challenge for individuals navigating the complexities of modern transportation systems. Traditional methods of disseminating such information, including static websites and printed materials, often lack interactivity and fail to provide tailored guidance to address specific queries and concerns. This deficiency leaves commuters vulnerable to misunderstandings, violations, and, in worst-case scenarios, accidents resulting from inadequate knowledge of traffic rules. Additionally, the sheer volume of regulations and updates makes it impractical for individuals to stay abreast of changes without a convenient and accessible means of accessing information.

Consequently, there exists a need for a dynamic, user-friendly, and technologically sophisticated solution capable of delivering real-time, contextually relevant guidance on traffic-related matters.

1.2 INTRODUCTION:

AI-based chatbots emerge as versatile instruments endowed with the capacity to furnish invaluable aid to users in navigating their queries and concerns. Whereas conventional modes of disseminating information regarding traffic regulations and laws often exhibit deficiencies in effectively engaging users, this project introduces a pioneering chat application poised to redefine the paradigm of traffic assistance services on a global scale.

Central to this endeavor is the integration of cutting-edge technologies spanning natural language processing, vector embeddings, and rigorous software engineering methodologies. Anchoring our methodology is the meticulous curation of Chroma, a vector database meticulously architected to house textual documents elucidating traffic

regulations, laws, and updates. By encoding this corpus of knowledge into vector embeddings, our chatbot attains a nuanced comprehension of contextual subtleties, thereby enabling the delivery of highly precise and contextually relevant responses to user inquiries.

At the helm of our chatbot's functionality is Mistral-7b-instruct, a meticulously calibrated and quantized language model. Exploiting contextual cues imprinted within the vector representations, Mistral-7b-instruct orchestrates responses that are not merely informative but also meticulously tailored to accommodate the specific needs and queries of users. To deftly bridge the chasm between vector embeddings and natural language, we leverage the capabilities of sentence transformers, which adeptly transmute embeddings into coherent and intelligible text, thereby facilitating seamless communication between users and the system.

Furthermore, our methodology encompasses a sophisticated retrieval-augmentation-generation (RAG) framework, wherein converted text undergoes iterative augmentation and subsequent response generation. This iterative refinement process ensures that the responses proffered by the chatbot are characterized by unparalleled accuracy, contextual resonance, and depth of information, thereby amplifying user satisfaction and engagement.

Augmenting the technological prowess underpinning our system is the robustness and versatility of the FastAPI framework. By seamlessly integrating with diverse platforms and applications while ensuring optimal efficiency and scalability, FastAPI fortifies the global applicability and accessibility of our AI-based chatbot solution. Through the symbiotic amalgamation of cutting-edge technologies and exacting software engineering methodologies, our project aspires to establish a new standard in the realm of AI-driven

traffic assistance, catalyzing safer and more informed commuting experiences for individuals across the globe.

CHAPTER - II

LITERATURE SURVEY

2.1 LITERATURE SURVEY

[1] Conversational information retrieval systems have garnered significant attention in recent years, reflecting the growing interest in enhancing user interactions with textual data. Smith et al. (2019) provided a comprehensive review of conversational information retrieval systems, emphasizing the need for natural language understanding and generation to improve user engagement. Advancements in language models, discussed by Chen and Manning (2020), particularly models like GPT-3, underscore the necessity for deeper contextual understanding and coherence in conversational agents. Deep learning techniques for document understanding tasks, as outlined by Zhang et al. (2021), showcase the potential of leveraging models like GPT-3 in improving document retrieval and comprehension. Interactive document retrieval systems, explored by Wang et al. (2022), highlight the importance of user engagement and relevance feedback, aligning with the interactive nature of the proposed "Chat with Documents using LLM" approach. Furthermore, Li et al. (2023) investigate integrating document understanding capabilities into conversational agents, enhancing user interactions by extracting relevant information from documents and incorporating it into dialogue contexts. By drawing from these studies, the proposed approach revolutionizes information retrieval by enabling users to engage in natural conversations with language models, bridging the gap between human-like conversation and document search.

[2] The research paper explores the integration of artificial intelligence (AI) in manufacturing through the development of a chatbot designed to assist users in completing assembly tasks akin to those encountered in manufacturing settings, with a focus on assembling a Meccanoid robot across multiple stages. Central to the study is the challenge of discerning users' intents accurately, particularly in multi-step tasks where queries may share identical intents but require distinct responses. To address this, the authors propose two innovative methodologies: firstly, incorporating visual features alongside textual features using the YOLO-based Masker with CNN (YMC) model to enable the chatbot to leverage visual cues for improved understanding and response generation; secondly, employing an Autoencoder to encode multi-modal features for enhanced intent classification, facilitating more accurate and relevant responses. Experimental results demonstrate significant improvements in the chatbot's performance following the integration of visual features, underscoring the efficacy of the proposed methodologies in capturing users' needs and providing tailored assistance, thus advancing the capabilities of AI-driven systems in manufacturing environments.

[3] The paper presents a pioneering contribution to the field of multimodal AI systems, aiming to enhance human-computer interaction through open-domain dialogues enriched with relevant photos. It addresses the limitations of existing approaches, particularly the lack of coherence between textual and visual inputs, by proposing a complete chatbot system consisting of two multimodal deep learning models: an image retriever utilizing ViT and BERT, and a response generator employing ViT and GPT-2/DialoGPT. These models are trained and evaluated on the PhotoChat dataset, showcasing superior performance compared to baseline systems in terms of image retrieval accuracy and response generation quality. The integration of image understanding and retrieval with text generation enables more engaging and contextually relevant conversations, as demonstrated

by human evaluation studies indicating higher image-groundedness, engagingness, fluency, coherence, and competitive humanness. Overall, the proposed multimodal chatbot system represents a significant advancement in the field, promising to revolutionize human-computer interaction by facilitating seamless and informative interactions between users and AI agents.

[4] A multi-modal chatbot seq2seq framework aimed at addressing the increased prevalence of mental illnesses, such as anxiety and depression, among young people exacerbated by the COVID-19 pandemic. With the shift to online learning and increased isolation from traditional social interactions, young individuals face heightened psychological challenges. The proposed model integrates text and image inputs from users to categorise the mental state of young individuals and offers a nuanced approach to understanding their needs. By combining image description and text summarization modules with an attention mechanism, the model effectively manages related content across different modalities. Experimental results on multi-modal datasets demonstrate a promising 70% average accuracy, while real user feedback confirms the system's efficacy in judgement. The paper highlights the importance of online psychotherapeutic applications, particularly amidst the pandemic's strain on mental health services. It underscores the necessity of remote interventions and the role of AI-driven chatbots in providing accessible mental health support. Additionally, it discusses the impact of the pandemic on various demographics, emphasising the unique challenges faced by adolescents and the importance of tailored interventions. Furthermore, the paper reviews related research on adolescent mental health, including studies on environmental influences, internet usage, and neurobehavioral changes during adolescence. It also examines existing chatbot applications in the mental health domain, showcasing their potential in providing emotional support and cognitive therapy. Overall, the paper contributes to the growing body of literature on leveraging AI

technologies to address mental health challenges, particularly among young populations navigating the complexities of the pandemic era.

[5] In the contemporary landscape, chatbots have become indispensable tools across diverse scientific disciplines. This research undertakes a comprehensive exploration into their utility, specifically honing in on sentence classification leveraging the News Aggregator Dataset as a litmus test to evaluate the model's performance vis-à-vis predetermined categories, thereby culminating in the creation of a robust chatbot program. Employing a multimodal approach, the study meticulously assesses four distinct models—namely GRU, Bi-GRU, 1D CNN, and 1D CNN Transpose—across six parameter variations to discern the most optimal configurations throughout the trial. Noteworthy is the revelation that the 1D CNN Transpose model emerges as the pinnacle performer, boasting an exceptional accuracy score of 0.9919. The research anticipates that both incarnations of the chatbot developed will not only furnish precise sentence predictions but also deliver discerning and accurate detection outcomes. Furthermore, the research meticulously elucidates each stage of the program's development, aiming to equip users with a comprehensive understanding that transcends mere operational proficiency, delving into the nuanced intricacies inherent within each sub-topic delineated within the study.

[6] A comprehensive framework aimed at quantitatively evaluating the capabilities of interactive Large Language Models (LLMs), particularly focusing on ChatGPT, by leveraging an extensive range of publicly available datasets spanning 23 datasets across eight diverse Natural Language Processing (NLP) application tasks. Through meticulous evaluation, encompassing multitask, multilingual, and multi-modal dimensions, the study provides a thorough analysis of ChatGPT's performance, including its ability to comprehend non-Latin script languages and generate multimodal content from textual prompts via an

intermediate code generation mechanism. Notably, the findings reveal ChatGPT's consistent outperformance of LLMs with zero-shot learning on a majority of tasks, and in some instances, even surpassing fine-tuned models. However, despite its proficiency in certain areas, such as language understanding, the study highlights ChatGPT's limitations, particularly in logical, non-textual, and commonsense reasoning tasks, where it exhibits an average accuracy of 63.41%, indicating its unreliability as a reasoner. Moreover, like other LLMs, ChatGPT faces challenges related to hallucination, further emphasizing the need for cautious interpretation of its outputs. Nevertheless, the paper underscores the interactive nature of ChatGPT, which enables collaboration with humans to enhance its performance through a multi-turn "prompt engineering" approach, leading to notable improvements in tasks such as summarization and machine translation. By releasing code for the extraction of evaluation sets, the paper facilitates the replication and extension of its findings, fostering ongoing research in the realm of interactive LLMs. [7] MMCHAT, a large-scale multi-modal dialogue corpus consisting of 32.4 million raw dialogues and 120.84 thousand filtered dialogues, aimed at integrating multi-modal contexts into conversation to enhance dialogue systems' engagement. Unlike previous corpora sourced from crowds or fictitious movies, MMCHAT comprises image-grounded dialogues extracted from real conversations on social media, where sparsity issues are observed. Notably, the dialogues often transition from image-initiated topics to non-image-grounded discussions over the course of the conversation. To address this challenge in dialogue generation tasks, the paper introduces a benchmark model incorporating an attention routing mechanism on image features. Experimental results underscore the utility of integrating image features and the model's effectiveness in mitigating the sparsity of such features, thereby advancing the capabilities of multi-modal dialogue systems.

[8] MultiModal Large Language Models (MM-LLMs) have witnessed significant progress, evolving to augment off-the-shelf LLMs and accommodate MultiModal (MM) inputs or outputs through cost-effective training methodologies. These advancements not only preserve the inherent reasoning and decision-making capabilities of LLMs but also extend their utility to a diverse array of MM tasks. This paper offers a comprehensive survey aimed at fostering further exploration and development within the MM-LLMs domain. Initially, the paper delineates general design formulations for model architecture and training pipelines. Subsequently, it presents a taxonomy comprising 122 MM-LLMs, each characterized by its distinct formulations. Furthermore, the survey assesses the performance of selected MM-LLMs on mainstream benchmarks and consolidates key training methodologies to bolster the effectiveness of MM-LLMs. Finally, the paper delves into prospective avenues for MM-LLMs advancement while concurrently providing a real-time tracking website for staying abreast of the latest developments in the field. It is hoped that this survey will significantly contribute to the ongoing progress and refinement of MM-LLMs, paving the way for future advancements in this burgeoning field.

[9] in-depth exploration of the burgeoning field of Multimodal Large Language Models (MLLMs), which represent a significant advancement in artificial intelligence research. MLLMs leverage the capabilities of Large Language Models (LLMs) to process and understand multimodal data, including text and images, enabling them to perform tasks that were previously challenging for traditional models. The survey begins by introducing the concept of MLLMs and outlining their fundamental principles. It then proceeds to discuss key techniques and applications employed in MLLM research, such as Multimodal Instruction Tuning (M-IT), Multimodal In-Context Learning (M-ICL), Multimodal Chain of Thought (M-CoT), and LLM-Aided Visual Reasoning (LAVR). Additionally, the survey identifies and addresses various challenges faced by researchers in this field, including data

scarcity and model complexity. Furthermore, it highlights promising research directions for future exploration and innovation in MLLMs. The survey concludes by emphasizing the importance of continuous updates and advancements in this rapidly evolving field, aiming to inspire further research and development in MLLMs.

CHAPTER - III

SYSTEM ANALYSIS

3.1 Methodology / Proposed Work

The proposed work for this project is to develop a comprehensive traffic guidance chatbot application that leverages advanced natural language processing techniques and vector databases to provide users with accurate, contextual, and informative responses regarding traffic rules, laws, and updates. The proposed work diagram given below




Fig. 3.1 Proposed Work Diagram

The proposed work aims to create a cutting-edge traffic guidance chatbot application that combines advanced natural language processing techniques, vector databases, and robust software engineering practices to provide users with a conversational and intuitive interface for accessing traffic-related information. By integrating retrieval, augmentation, and generation processes, the system ensures contextually relevant and informative responses, contributing to safer and more informed commuting experiences.

The key components and steps involved in the proposed work are as follows:

1. Data Collection and Vectorization:

- Conduct an extensive research to identify authoritative sources for traffic rules, regulations, and updates, including government websites, legal documents, official publications, and credible online resources.
- Gather a comprehensive corpus of textual data from these sources, ensuring the inclusion of up-to-date information and covering a wide range of topics related to traffic guidance.
- Preprocess the collected data by performing cleaning operations such as removing irrelevant content, formatting inconsistencies, and handling specific data formats (e.g., PDF, scanned documents).
- Utilize state-of-the-art sentence transformer model **BAAI/bge-large-en-v1.5**, which have been pre-trained on large corpora, to convert the preprocessed textual data into dense vector representations (embeddings).
- Optional- fine-tune the large language model **mistralai/Mistral-7B-Instruct-v0.2** on the collected traffic data to obtain domain-specific embeddings tailored for the task.

- Store the resulting vector embeddings in a vector database called Chroma, which uses efficient data structures and indexing techniques for fast retrieval and similarity search operations.

2. Language Model Fine-Tuning and Quantization:

- Select a large pre-trained language model like Mistral-7b-instruct, which has been trained on a vast amount of text data and has demonstrated strong performance in natural language understanding and generation tasks.
- Fine-tune the selected language model on the domain-specific data (traffic rules and regulations) using techniques like transfer learning or continued pre-training.
- During fine-tuning, expose the model to a diverse set of traffic-related queries and their corresponding responses to improve its understanding of the domain and its ability to generate accurate and relevant responses.
- Employ quantization techniques like post-training quantization or quantization-aware training to reduce the size and computational requirements of the fine-tuned language model.
- Quantization can involve techniques like weight pruning, weight sharing, or low-precision arithmetic operations, while maintaining the model's performance and accuracy.

3. Retrieval-Augmentation-Generation (RAG) Process:

- Implement a retrieval mechanism that takes a user's query as input and retrieves the most relevant vector embeddings from the Chroma database using techniques like nearest neighbor search or maximum inner product search.

- Develop an augmentation component that converts the retrieved vector embeddings back into readable text snippets using sentence transformer models or other techniques like nearest neighbor search in the embedding space.
- Integrate the fine-tuned and quantized language model with the retrieved and augmented text snippets to generate a final response to the user's query.
- Employ techniques like attention mechanisms or cross-attention to allow the language model to effectively leverage the context provided by the retrieved and augmented information.
- Explore different architectures and approaches for the RAG process, such as end-to-end models or modular architectures, to optimize performance and accuracy.

4. User Interface and Integration:

- Design and develop a user-friendly interface, either web-based or mobile, that allows users to interact with the chatbot through chat sessions seamlessly.
- Implement features like input validation, error handling, and session management to ensure a smooth user experience.
- Leverage the FastAPI framework to build the backend API that integrates the vector database, language model, and the RAG process.
- Implement caching mechanisms and asynchronous processing to improve the system's responsiveness and handle concurrent user requests efficiently.
- Ensure secure communication between the frontend and backend components by implementing authentication and authorization mechanisms, as well as data encryption and protection measures.

5. Evaluation and Testing:

- Develop a comprehensive test suite that covers various aspects of the system, including unit tests, integration tests, and end-to-end tests.

- Define evaluation metrics to assess the chatbot's performance, such as accuracy, relevance, informativeness, and user satisfaction.
- Conduct automated testing and manual evaluation by domain experts to ensure the correctness and quality of the generated responses.
- Perform load testing and stress testing to evaluate the system's performance under different load conditions and identify potential bottlenecks or scalability issues.
- Continuously monitor the system's performance in production and gather user feedback to identify areas for improvement and iterate on the system's design and implementation.
- Conduct user studies and usability testing to evaluate the chatbot's effectiveness and identify opportunities for enhancing the user experience.

6. Continuous Improvement and Updates:

- Establish a process for regularly monitoring and incorporating updates to traffic rules, laws, and regulations from authoritative sources.
- Implement a mechanism to update the vector database with new or modified information, ensuring that the system remains up-to-date and accurate.
- Periodically fine-tune the language model on the updated data to maintain and improve its performance and understanding of the domain.
- Gather user feedback and analytics to identify areas for improvement, such as frequently asked questions, common misunderstandings, or gaps in the system's knowledge.

By following this detailed proposed work, the project aims to create a state-of-the-art traffic guidance chatbot application that leverages advanced natural language processing techniques, vector databases, and robust software engineering practices. The system will provide users with accurate, contextual, and informative responses

regarding traffic rules, laws, and updates, ultimately contributing to safer and more informed commuting experiences

3.2 SYSTEM ANALYSIS

The existing landscape of accessing comprehensive and current information on traffic regulations and updates highlights a significant challenge for individuals navigating modern transportation systems. Conventional methods like static websites and printed materials lack interactivity and fail to provide tailored guidance, leaving commuters vulnerable to misunderstandings and violations. Additionally, the volume of regulations makes it impractical for individuals to stay updated without a convenient means of accessing information.

To address this challenge, a dynamic and user-friendly solution is proposed, leveraging AI-based chatbots to deliver real-time, contextually relevant guidance on traffic-related matters. This project aims to redefine traffic assistance services globally through the integration of cutting-edge technologies such as natural language processing (NLP) and vector embeddings.

At the core of this solution is Chroma, a meticulously curated vector database housing textual documents elucidating traffic regulations. By encoding this knowledge into vector embeddings, the chatbot achieves a nuanced understanding of contextual subtleties, enabling precise and relevant responses to user inquiries. The chatbot's functionality is powered by Mistral-7b-instruct, a calibrated language model adept at orchestrating tailored responses based on contextual cues within the vector representations.

To bridge the gap between vector embeddings and natural language, sentence transformers are employed to transmute embeddings into coherent text, facilitating

seamless communication. The methodology also includes a sophisticated retrieval-augmentation-generation (RAG) framework, ensuring that responses are accurate, contextually resonant, and informative through iterative refinement.

The system's technological prowess is augmented by the robustness and versatility of the FastAPI framework, enabling seamless integration with diverse platforms and applications. This integration enhances the global applicability and accessibility of the AI-based chatbot solution, aiming to establish a new standard in AI-driven traffic assistance for safer and more informed commuting experiences worldwide.

3.3 EXISTING SYSTEM

3.3.1 Rule-based Chatbot

Chatbots with Rule-based Systems: Traditional chatbots relied on rule-based systems to respond to user queries and commands. These systems operated based on predefined rules and patterns, lacking the ability to adapt or learn from user interactions dynamically. While they could handle basic inquiries, their responses were often rigid and lacked the sophistication of AI-driven approaches.

The existing system for developing chatbots to collect user self-reported data and facilitate natural language conversations encounters several limitations, necessitating a shift towards more sophisticated approaches. Traditional chatbots, often rule-based or script-driven, struggle to engage in dynamic, flexible dialogues and provide personalized responses tailored to individual users. Their reliance on predefined rules and scripts can lead to rigid and unnatural interactions, limiting their effectiveness in capturing the nuances of human language and context.

3.3.2 Leveraging Large Language Models for Chatbots

Recent advancements in natural language processing (NLP) and the emergence of large language models (LLMs) present an opportunity to overcome these limitations and develop more human-like chatbots. Models like GPT-3 and PaLM showcase remarkable language understanding and generation capabilities, enabling them to engage in nuanced, contextually relevant conversations.

However, leveraging LLMs for task-oriented chatbots, particularly in domains like collecting user self-reports, presents its own challenges. Designing effective prompts and fine-tuning strategies for LLMs to accurately interpret user inputs and provide contextually relevant responses requires extensive research and experimentation. Furthermore, the computational complexity and resource demands of existing LLMs pose challenges for deployment on resource-constrained devices or environments.

While efforts have been made to develop more efficient LLMs, such as LLaMA and Baize, their performance and suitability for specific use cases like collecting user self-reports require further evaluation. Additionally, concerns persist regarding the interpretability and transparency of LLM-based chatbots, particularly in sensitive domains like healthcare or finance, where understanding the reasoning behind their responses is crucial.

LLMs hold significant promise for developing more natural and engaging chatbots, the existing system lacks a comprehensive solution that addresses the challenges of prompt design, task-specific fine-tuning, resource efficiency, and interpretability, particularly in the context of collecting user self-reported data through natural language conversations.

3.4 Social Feasibility

- **User Acceptance:** The success of TrafficGPT hinges on user acceptance and adoption. Understanding whether the target audience, such as commuters and drivers, will embrace its features and functionalities is pivotal. Factors like preferences for intuitive traffic guidance, concerns regarding data privacy, and the perceived usefulness of AI-driven assistance will heavily influence user acceptance.
- **Privacy and Security Considerations:** In an age where data privacy and security are paramount concerns, users are increasingly cautious about sharing personal information online. TrafficGPT's commitment to prioritizing secure data handling and processing addresses these concerns, potentially enhancing user trust and confidence in the platform. Ensuring that users perceive the application as a secure and reliable source of traffic information is crucial for its social feasibility.
- **Inclusivity and Accessibility:** Social feasibility also involves considerations of inclusivity and accessibility. TrafficGPT's compatibility across various devices and its user-friendly interface aim to make traffic guidance accessible to a wide range of users, regardless of their technical expertise or device preferences. By catering to diverse user needs, TrafficGPT strives to enhance accessibility and inclusivity in the realm of traffic assistance services.
- **Educational and Professional Impact:** TrafficGPT's focus on providing accurate and timely traffic information can have positive social implications. By empowering users with knowledge about traffic rules and updates, TrafficGPT contributes to safer and more informed commuting experiences. Additionally, its educational resources and real-time updates can aid in professional development

for individuals in transportation-related fields, aligning with broader societal goals of safety and efficiency in transportation systems.

3.5 HARDWARE AND SOFTWARE REQUIREMENTS

Our chatbot system relies on powerful hardware and software components to deliver efficient and intuitive user experiences. By leveraging advanced technologies such as PyTorch, Hugging Face Transformers, ChromaDB, and Streamlit, coupled with robust hardware configurations, we aim to create a high-performing chatbot solution capable of seamless AI integration, efficient database management, and intuitive user interface development.

3.5.1 Hardware Requirements

- **Processor:** Intel Core i5 processor or higher for optimal performance.
- **RAM:** Minimum of 16GB RAM for smooth operation, although higher RAM configurations are recommended for better performance.
- **Storage:** SSD storage with a capacity of at least 256GB for faster read/write speeds and improved overall system responsiveness.
- **GPU:** NVIDIA GeForce or AMD Radeon GPU for accelerated AI model

inference, enhancing the speed and efficiency of processing.

3.5.2 Software Requirements:

- **AI Model Integration:** Utilize PyTorch and Hugging Face Transformers libraries for seamless integration of advanced AI models, ensuring robust natural language processing capabilities.
- **Database Management:** Employ SQLite or ChromaDB for lightweight and

efficient database management, facilitating the storage of chat history and user data securely.

- **User Interface Development:** Implemented HTML, CSS, and JavaScript for

crafting the user interface to ensure a visually appealing and responsive design. This approach enhances user experience through intuitive interactions. Integrate the frontend with the backend using FastAPI to seamlessly combine both ends of the application.

CHAPTER - IV

SYSTEM DESIGN

4.1 FLOW DIAGRAM :

The flow diagram deals with the implementation design of the model. It defines the steps that are involved since the start of the process. The flow diagram is the visual implementation of the whole process that is carried throughout the entire process. The flow diagram is presented below as Figure 4.1




Fig No 4.1: Flow Diagram

CHAPTER - V

SYSTEM ARCHITECTURE

5.1 SYSTEM ARCHITECTURE

A system architecture is a conceptual model that defines the structure, behavior, and views of a system. An architecture description is a formal description and representation of a system, organized in a way that supports reasoning about the structures and behaviors of the system. The system architecture diagram, as depicted in Figure 5.1, is presented below.




Fig No 5.1 System Architecture

The architectural diagram of TrafficGPT delineates a sophisticated process by which input data undergoes transformation into embeddings, which are then stored within vector databases. This initial stage involves the conversion of raw input data, such as text, documents, or any other form of information, into numerical representations known as embeddings. These embeddings serve as condensed, high-dimensional vectors that encapsulate the semantic meaning and contextual information of the input data. Following this conversion, the embeddings are systematically organized and stored within vector databases, allowing for efficient retrieval and manipulation.

Subsequently, when prompted with inquiries or tasks, the system retrieves relevant embeddings from the vector databases. This retrieval process is guided by the nature of the

prompt, which involves natural language queries as input. The retrieved embeddings are then utilized in various computational processes to generate responses, perform analyses, or execute specific actions as required by the given task.

Throughout this intricate architecture, the embeddings serve as the fundamental units of information exchange, enabling seamless communication between input data and the system's computational capabilities.

5.2 Modules:

Modules are designed to be modular, meaning they can be developed, tested, and maintained independently, and can interact with other modules through well-defined interfaces.

5.2.1 Module 1: Data Loading and Embedding conversion

Module 1 establishes the foundational concepts of data loading and embedding. The process begins with uploading a data file to the "docs" folder within the project directory. This data file must be in PDF format. The figure for Module 1 is presented below as Figure 5.2.1.




Fig No 5.2.1 - Data Loading and Embedding conversion

Step 1: Data Upload

The initial step of the data processing workflow involves the user uploading files to a designated folder named "docs" within the project directory. These files are typically in PDF format and contain the textual content that will undergo further processing.

Step 2: Execution of PDF Ingestion Script

Following the data upload, the user navigates to the project directory via the command prompt interface. Hereafter, the user invokes the execution of a Python script named "pdf_ingest.py" by issuing the command "python pdf_ingest.py". It is imperative to note that the command prompt must be positioned within the project folder for seamless execution.

Step 3: Document Segmentation

Upon execution of the "pdf_ingest.py" script, the uploaded documents undergo segmentation into smaller, more manageable units or "chunks". This segmentation process ensures that the subsequent embedding generation operates efficiently by handling smaller portions of the document at a time.

Step 4: Embedding Generation

Following document segmentation, each chunk undergoes embedding generation using the "bge-large-en-v1.5" model. This model employs advanced Natural Language Processing techniques to transform textual data into dense vector representations, effectively capturing semantic relationships and contextual information embedded within the text.

Step 5: Storage in Choma Vector Database

The resultant embeddings are then stored in the Choma vector database, a purpose-built database optimized for the storage and retrieval of vectorized data. Storing embeddings in this database enables rapid and efficient querying, facilitating the retrieval of relevant information in response to user queries or search requests.

This comprehensive data processing workflow seamlessly integrates document ingestion, segmentation, embedding generation, and storage within a structured framework. By leveraging advanced techniques in Natural Language Processing and efficient database storage, the workflow ensures optimal handling of textual data, empowering users with the ability to efficiently retrieve information and extract insights from large document repositories.

Module 2: Querying from Loaded Data

Module 2, introduces the fundamental principles of data querying within the context of vector databases. Vector databases offer a powerful mechanism for storing and retrieving information. This module will explore the core functionalities involved in querying loaded data, providing a foundation for further exploration of this essential skill. The figure for Module 2 is presented below as Figure 5.2.2.




Fig No. 5.2.2 : Querying from Loaded Data

Step 1: User Interaction:

Users engage with the TrafficGPT system through a sophisticated user interface crafted with HTML, CSS, and JavaScript. The interface provides a seamless platform for users to submit prompts and queries relevant to their needs. Rather than relying on conventional Streamlit, TrafficGPT leverages custom-built web technologies to deliver an intuitive and responsive user experience. FastAPI serves as the bridge, seamlessly integrating the frontend and backend components of the system.

Step 2: Prompt Processing:

Upon receipt of a user prompt, TrafficGPT initiates a comprehensive processing pipeline to extract actionable insights. The textual prompt undergoes transformation via the bge-large-en-v1.5 language model, converting it into dense embeddings. These embeddings encapsulate the semantic essence and contextual nuances of the user's query, laying the groundwork for efficient search and retrieval operations.

Step 3: Similarity Search:

The generated embeddings are then utilized to conduct a similarity search within the ChromaDB. This database houses embeddings extracted from previously processed text chunks sourced from various documents. The objective of this search operation is to pinpoint text chunks closely aligned with the semantic content of the user's prompt. Leveraging advanced similarity metrics, TrafficGPT retrieves relevant text chunks based on their proximity to the prompt embeddings.

Step 4: Combination and Processing:

Subsequently, the user prompt and the retrieved text chunks are amalgamated and inputted into the Mistral-7b-Instruct-v0.2-GGUF model. Renowned for its adeptness in contextual understanding and semantic inference, this model undertakes comprehensive processing of the combined input. By discerning context and meaning, Mistral-7b-Instruct-v0.2-GGUF generates a coherent and tailored response tailored precisely to the user's query.

Step 5: Response Storage:

Upon completion of response generation, both the original prompt and the generated response are archived as part of the chat session. This ensures a cohesive user experience, allowing users to access previous interactions and responses seamlessly. In the web application interface, previous responses are prominently displayed, enriching the user experience and fostering continuity in dialogue.

CHAPTER VI

SYSTEM IMPLEMENTATION

6. System Implementation

This chapter describes the implementation details of the project. It describes the different modules of the proposed system, the tools and platforms used. Implementation details are also mentioned and output snapshots for each module are shown. The following presents the input and output models used in the implementation of the proposed system

6.1 System Implementation Requirements

The following software dependencies are required for the system implementation:

- aiohttp (Version 3.9.3)
- aiosignal (Version 1.3.1)
- annotated-types (Version 0.6.0)
- antlr4-python3-runtime (Version 4.9.3)
- anyio (Version 4.3.0)
- async-timeout (Version 4.0.3)
- attrs (Version 23.2.0)
- backoff (Version 2.2.1)
- beautifulsoup4 (Version 4.12.3)
- certifi (Version 2024.2.2)
- cffi (Version 1.16.0)
- chardet (Version 5.2.0)
- charset-normalizer (Version 3.3.2)
- click (Version 8.1.7)
- colorama (Version 0.4.6)
- coloredlogs (Version 15.0.1)
- contourpy (Version 1.2.0)
- cryptography (Version 42.0.4)
- ctransformers (Version 0.2.27)
- cyclcr (Version 0.12.1)
- dataclasses-json (Version 0.6.4)

- dataclasses-json-speakeasy (Version 0.5.11)
- Deprecated (Version 1.2.14)
- diskcache (Version 5.6.3)
- effdet (Version 0.4.1)
- emoji (Version 2.10.1)
- et-xmlfile (Version 1.1.0)
- exceptiongroup (Version 1.2.0)
- fastapi
- filelock (Version 3.13.1)
- filetype (Version 1.2.0)

-
-
-
-
-
-
-
- flatbuffers (Version 23.5.26)
- fonttools (Version 4.49.0)
- frozenset (Version 1.4.1)
- fsspec (Version 2024.2.0)
- greenlet (Version 3.0.3)
- grpcio (Version 1.60.1)
- grpcio-tools (Version 1.60.1)
- h11 (Version 0.14.0)
- h2 (Version 4.1.0)
- hpack (Version 4.0.0)
- httpcore (Version 1.0.3)
- httpx (Version 0.26.0)
- huggingface-hub (Version 0.20.3)
- humanfriendly (Version 10.0)
- hyperframe (Version 6.0.1)
- idna (Version 3.6)
- iopath (Version 0.1.10)
- Jinja2 (Version 3.1.3) • joblib (Version 1.3.2)
- jsonpatch (Version 1.33)
- jsonpath-python (Version 1.0.6)
- jsonpointer (Version 2.4)
- kiwisolver (Version 1.4.5)
- langchain (Version 0.1.8)
- langchain-community (Version 0.0.21)
- langchain-core (Version 0.1.25)
- langdetect (Version 1.0.9)
- langsmith (Version 0.1.5)
- layoutparser (Version 0.3.4)
- llama_cpp_python (Version 0.2.44)
- lxml (Version 5.1.0)
- Markdown (Version 3.5.2)
- MarkupSafe (Version 2.1.5)

-
-
-
-
-
-
-
-
- marshmallow (Version 3.20.2)
- matplotlib (Version 3.8.3)
- mpmath (Version 1.3.0)
- msg-parser (Version 1.2.0)
- multidict (Version 6.0.5)
- mypy-extensions (Version 1.0.0)
- networkx (Version 3.2.1)
- nltk (Version 3.8.1)
- numpy (Version 1.26.4)
- olefile (Version 0.47)
- omegaconf (Version 2.3.0)
- onnx (Version 1.15.0)
- onnxruntime (Version 1.15.1)
- opencv-python (Version 4.9.0.80) openpyxl (Version 3.1.2) packaging (Version 23.2)
- pandas (Version 2.2.0)
- pdf2image (Version 1.17.0)
- pdfminer.six (Version 20221105)
- pdfplumber (Version 0.10.4)
- pikepdf (Version 8.13.0)
- Pillow (Version 10.2.0)
- Pillow_Heif (Version 0.15.0)
- portalocker (Version 2.8.2)
- protobuf (Version 4.25.3)
- Py-Cpuinfo (Version 9.0.0)
- PyCocotools (Version 2.0.7)
- PyParser (Version 2.21)
- Pydantic (Version 2.6.1)
- Pydantic_Core (Version 2.16.2)
- PyPandoc (Version 1.13)
- PyParsing (Version 3.1.1)

-
-
-
-
-
-
-
-
- pypdf (Version 4.0.2)
- pypdfium2 (Version 4.27.0)
- pyreadline3 (Version 3.4.1)
- pytesseract (Version 0.3.10)
- python-dateutil (Version 2.8.2)
- python-docx (Version 1.1.0)
- python-iso639 (Version 2024.2.7)
- python-magic (Version 0.4.27)
- python-multipart (Version 0.0.9)
- python-pptx (Version 0.6.23)
- pytz (Version 2024.1)
- pywin32 (Version 306)
- PyYAML (Version 6.0.1)
- qdrant-client (Version 1.7.3)
- rapidfuzz (Version 3.6.1)
- regex (Version 2023.12.25)
- requests (Version 2.31.0)
- safetensors (Version 0.4.2)
- scikit-learn (Version 1.4.1.post1)
- scipy (Version 1.12.0)
- sentence-transformers (Version 2.3.1)
- sentencepiece (Version 0.2.0)
- six (Version 1.16.0)
- sniffio (Version 1.3.0)
- soupsieve (Version 2.5)
- SQLAlchemy (Version 2.0.27)
- starlette (Version 0.36.3)
- sympy (Version 1.12)
- tabulate (Version 0.9.0)
- tenacity (Version 8.2.3)
- threadpoolctl (Version 3.3.0)
- timm (Version 0.9.16)

-
-
-
-
-
-
-
-
- tokenizers (Version 0.15.2)
- torch (Version 2.2.0)
- torchvision (Version 0.17.0)
- tqdm (Version 4.66.2)
- transformers (Version 4.37.2)
- typing-inspect (Version 0.9.0)
- typing_extensions (Version 4.9.0)
- tzdata (Version 2024.1)
- unstructured
- unstructured-client
- unstructured-inference
- unstructured.pytesseract
- urllib3 (Version 2.2.1)
- uvicorn (Version 0.27.1)
- wrapt (Version 1.16.0)
- xlrd (Version 2.0.1)
- XlsxWriter (Version 3.2.0)
- yarl (Version 1.9.4)

6.2 Python Implementation of TrafficGPT (Backend code)

```
# Import necessary modules and classes from
langchain.prompts import PromptTemplate from
langchain_community.llms import LlamaCpp
from langchain.chains import RetrievalQA
from langchain_community.embeddings import
SentenceTransformerEmbeddings from fastapi import FastAPI, Request, Form,
Response from fastapi.responses import HTMLResponse from fastapi.templating
import Jinja2Templates from fastapi.staticfiles import StaticFiles from
fastapi.encoders import jsonable_encoder from
langchain_community.vectorstores import Chroma import os import json
from langchain.callbacks.manager import CallbackManager
```

```

from langchain.callbacks.streaming_stdout import
StreamingStdOutCallbackHandler

# Initialize FastAPI app
app = FastAPI()

# Initialize CallbackManager for handling callbacks callback_manager =
CallbackManager([StreamingStdOutCallbackHandler()])

# Initialize Jinja2Templates for rendering HTML templates
templates = Jinja2Templates(directory="templates")

# Mount static files directory app.mount("/static",
StaticFiles(directory="static"), name="static")

# Define local LLM model path
local_llm = "models/mistral-7b-instruct-v0.2.Q5_K_M.gguf"

# Initialize LlamaCpp
model llm = LlamaCpp(
    model_path=local_llm,
    temperature=0.1,
    max_tokens=5000, top_p=1,
    n_ctx=2048, verbose=True,
    device='cuda',
    callback_manager=callback_manager,
    n_gpu_layers=1,
    n_batch=2048
)

```

Explanation of the above arguments:

llamacpp: LLaMA is a powerful large language model developed by Meta AI, and LLaMaCpp is a C++ implementation that allows you to run the LLaMA model efficiently on various hardware platforms, including GPUs.

1. **model_path=local_llm:** This argument specifies the path to the pre-trained LLaMA model. The value local_llm is likely a variable holding the path to the

model file on your local machine. It is essential to provide the correct path to the model file for the LLaMA model to load correctly.

2. **temperature=0.1:** The temperature parameter controls the randomness of the generated text. A lower temperature value (e.g., 0.1) makes the output more focused and deterministic, while a higher temperature value (e.g., 0.8) makes the output more diverse and unpredictable. The optimal value depends on your use case. For tasks that require precise and accurate output, a lower temperature (e.g., 0.1 or lower) is recommended. For creative writing or exploration, a higher temperature (e.g., 0.5 to 0.8) might be more suitable.
3. **max_tokens=5000:** This argument specifies the maximum number of tokens (words or subwords) that the model can generate. The optimal value depends on your use case and available computational resources. A higher value allows for longer output but requires more memory and computation time. A value of 5000 is generally suitable for most tasks, but you may need to adjust it based on your specific requirements.
4. **top_p=1:** The top_p parameter controls the nucleus sampling strategy used for text generation. A value of 1 means that the model will consider all possible tokens for generation. Lower values (e.g., 0.9 or 0.8) restrict the model to only consider the most probable tokens, potentially improving the quality of the output but also reducing diversity. The optimal value depends on your use case. For most tasks, a value of 1 is recommended, but you may want to experiment with lower values if you need more controlled or focused output.
5. **n_ctx=2048:** This argument specifies the maximum context length that the model can consider for generating the next token. The value of 2048 is a common choice for LLaMA models, as it allows for a reasonably long context while keeping the computational requirements manageable. The optimal value depends on your use case and available computational resources. For tasks that require longer context (e.g., summarizing long documents), you may need to increase this value, but doing so will also increase the memory and computational requirements.

6. **verbose=True:** This argument controls the verbosity of the LLaMA model's output. Setting it to True will print detailed information about the model's progress and performance, which can be useful for debugging or monitoring purposes. The optimal value depends on your preference and whether you need to see detailed output or not.
7. **device='cuda':** This argument specifies the device to use for computation. The value 'cuda' indicates that the model should run on a CUDA-enabled GPU, which can significantly speed up computations compared to running on a CPU. If you have a compatible GPU, using 'cuda' is recommended for optimal performance. If you don't have a GPU or want to run on the CPU, you can set this value to 'cpu'.
8. **callback_manager=callback_manager:** This argument allows you to pass a custom callback manager to the LLaMA model. Callback managers can be used to monitor and control the model's behavior during inference. The optimal value depends on your specific requirements and whether you need to implement custom callbacks or not.
9. **n_gpu_layers=1:** This argument specifies the number of GPU layers to use for model parallelism. The optimal value depends on the available GPU memory and the size of the LLaMA model. For smaller models or GPUs with limited memory, a value of 1 (no model parallelism) is recommended. For larger models or GPUs with more memory, you can increase this value to distribute the model across multiple GPU layers, which can improve performance but also increases complexity. Note: Mistral-7b-instruct model requires 33 layers to fit in the GPU."
10. **n_batch=2048:** The n_batch argument specifies the batch size for inference, determining how many input sequences the LLaMA model processes in parallel. Its optimal value depends on the available GPU memory and the model size. While larger batch sizes can enhance performance through parallelism, they must not exceed the GPU memory limits to avoid degradation or errors. Typically set to 2048 for LLaMA models, this value balances performance and memory usage but may require adjustment based on the specific hardware resources and model size. Experimentation is necessary to find the optimal batch size within the range of 1

to n_ctx (maximum context length), considering the amount of VRAM available on the GPU to maximize performance while respecting memory constraints.

Print message indicating LLM initialization

```
print("LLM Initialized....")
```

Define prompt template

```
prompt_template = """ Most of the time use  
the following pieces of information to  
answer the user's question.
```

Your name is TrafficGPT, developed by Greater Chennai Traffic Police (GCTP) to provide comprehensive guidance on traffic rules, laws, and updates of Chennai.

Context: {context}

Question: {question}

Answer must be perfect and well
explained. answer: """

```
# Initialize SentenceTransformerEmbeddings embeddings =  
SentenceTransformerEmbeddings(model_name="bge-large-en-v1.5")
```

Define persist directory for vector database

```
persist_directory = "vector database"
```

Initialize Chroma vector store

```
vectordb = Chroma(persist_directory=persist_directory,  
embedding_function=embeddings)
```

Define prompt template prompt =

```
PromptTemplate(template=prompt_template, input_variables=['context',  
'question'])
```

```
# Create retriever from Chroma vector store retriever  
= vectordb.as_retriever(search_kwargs={"k": 1})
```

Initialize conversation history dictionary

```
conversation_history = {}
```

Define route to serve HTML template

```
@app.get("/", response_class=HTMLResponse)
async def read_root(request: Request):
    formatted_history = "<br>".join([f"<strong>Query:</strong>
{query}<br><strong>Response:</strong> {response}" for query, response
in conversation_history.items()]) return
templates.TemplateResponse("index.html", {"request": request,
"conversation_history": formatted_history})
```

Define route to handle POST requests for getting responses

```
@app.post("/get_response") async def
get_response(query: str = Form(...)): # Define
chain type kwargs chain_type_kwargs =
{"prompt": prompt}
```

```
    # Initialize RetrievalQA model qa = RetrievalQA.from_chain_type(llm=llm,
chain_type="stuff", retriever=retriever,
chain_type_kwargs=chain_type_kwargs, verbose=True)
```

Get response from the QA model for the given query

```
response = qa(query)
```

Extract the answer from the response

```
answer = response['result']
```

Append the current query and response to conversation history

```
conversation_history[query] = answer
```

Encode response data as JSON

```
response_data = jsonable_encoder(json.dumps({"answer": answer}))
```

Create an HTTP response object

```
http_response = Response(response_data)
```

Return the HTTP response

```
return http_response
```

6.3 Knowledge Base Update Process:

1. Identify New and Updated Documents:

This step involves identifying any documents that have been added or modified since the last update of the knowledge base. It's essential to review these documents to determine if they contain information that should be incorporated into the chatbot's knowledge base.

2. Delete Old Embedding Database:

Before updating the knowledge base, it's crucial to remove the existing embedding database entirely. This ensures that there are no remnants of outdated information, providing a clean slate for the update process.

3. Prepare Document Data:

In this step, documents are prepared for processing by the vectorization system. This may involve converting documents to a format that can be efficiently handled by the system, such as plain text. Ensuring that documents are in the appropriate format is essential for accurate vectorization and embedding generation.

4. Update Vector Database:

Utilizing ChromaDB, the vector database is updated with embeddings for the new and updated documents. Only embeddings for documents present in the 'docs' folder are added, maintaining the relevancy and integrity of the knowledge base. This step ensures that the chatbot has access to the latest information when generating responses to user queries.

5. Verify and Test:


After updating the knowledge base, it's essential to verify the accuracy and functionality of the updated information. This involves testing the chatbot's responses against various queries to ensure that it correctly utilizes the newly added or modified information.

Verification and testing help identify any potential issues or inaccuracies that may need to be addressed before deploying the updated knowledge base.

6. Document the Changes:

Finally, it's crucial to document all modifications made to the knowledge base during the update process. This documentation serves as a record of changes for future reference and troubleshooting purposes. It provides valuable insights into the evolution of the knowledge base over time, aiding in its maintenance and continuous improvement.

6.4 Output of the Chat Application with a prompt and Response



The provided image depicts a screenshot of a chat application interface linked with the Greater Chennai Traffic Police (GCTP). Through the interface, users engage in conversations, specifically related to vehicle registration certificates. The exchange suggests that the chat application functions as a user-friendly platform for individuals to inquire about or request new registration certificates for their vehicles. This convenience potentially streamlines the application and inquiry process, offering a more accessible avenue for users to interact with authorities regarding vehicle registration matters.

In the conversation snippet, the chatbot succinctly outlines the penalty associated with driving without a valid registration certificate: Rs. 5,000. This penalty is a significant deterrent, intended to encourage compliance with regulatory mandates regarding vehicle registration. By imposing fines, authorities seek to emphasize the importance of adhering to legal requirements, ensuring that vehicles are properly registered and meet safety standards before being driven on public roads. The specified penalty amount reflects the gravity of the offense and underscores the significance of regulatory compliance in maintaining road safety and order.

The stored messages are displayed using separate text sections or cards, with each card containing two distinct components. The first component is the "Query" section, which displays the previously asked question or query. For example, one of the queries shown is "Query: penalty for driving without Registration Certificate." This section clearly indicates the specific question or topic that was previously inquired about.

The second component is the "Response" section, which provides the answer or information related to the corresponding query. The response is presented in a concise and informative manner, explaining the relevant details, rules, and penalties associated with the query. For instance, the response to the aforementioned query outlines the

penalty amount and legal compliance requirements for driving without a valid registration certificate under the Motor Vehicle Act.

The separate cards for queries and responses creates a well-organized and easily readable presentation of the stored messages. Users can clearly distinguish between the asked questions and the provided answers, making it convenient to review and understand the information.

CHAPTER - VII

CONCLUSIONS AND FUTURE ENHANCEMENTS

7.1 Conclusion:

The proposed TrafficGPT project represents a significant advancement in the field of AI-powered traffic guidance systems. By combining state-of-the-art natural language processing techniques, vector databases, and robust software engineering practices, this project aims to revolutionize the accessibility and efficacy of traffic-related information dissemination.

Through the integration of a meticulously curated vector database (Chroma), a fine-tuned and quantized language model (Mistral-7b-instruct), and a retrieval-augmentation-generation (RAG) process, the TrafficGPT chatbot system provides users with accurate, contextual, and informative responses to their traffic-related queries. The system's ability to understand the nuances of natural language and leverage contextual cues from vector embeddings ensures that responses are tailored to the specific needs and concerns of users.

Furthermore, the utilization of the FastAPI framework enhances the system's efficiency, scalability, and seamless integration with various platforms and applications, making it accessible to a wide range of users across different devices and modalities.

By addressing the challenges of accessing up-to-date and comprehensive traffic information, the TrafficGPT project has the potential to significantly improve road safety, reduce traffic violations, and enhance the overall commuting experience for individuals navigating complex transportation systems.

7.2 Future Enhancements:

While the proposed TrafficGPT project represents a significant milestone in the development of AI-powered traffic guidance systems, there are several avenues for future research and development:

1. **Multilingual Support:** Extending the system to support multiple languages would greatly enhance its global reach and accessibility, catering to diverse linguistic communities worldwide.

2. **Multimodal Integration:** Incorporating multimodal inputs, such as images, videos, or voice commands, could further enhance the user experience and enable more natural and intuitive interactions with the chatbot.
3. **Personalization and Context Awareness:** Implementing mechanisms to learn and adapt to individual user preferences, locations, and commuting patterns could lead to more personalized and context-aware responses, improving the overall relevance and usefulness of the system.
4. **Integration with Real-time Traffic Data:** Combining the system with live traffic data sources, such as traffic cameras, sensor networks, and crowdsourced information, could enable real-time updates and guidance tailored to current traffic conditions.
5. **Proactive Guidance and Prediction:** Exploring predictive models to anticipate potential traffic issues, accidents, or congestion could enable the system to provide proactive guidance and recommendations, helping users plan their routes and travel times more effectively.

CHAPTER - VIII

REFERENCES

REFERENCES

- [1] P. Sujit, S. Saripalli, and J. B. Sousa, “Unmanned aerial vehicle path following: A survey and analysis of algorithms for fixed-wing unmanned aerial vehicles,” *IEEE Control Syst. Mag.*, vol. 34, no. 1, pp. 42–59, Feb. 2014.
- [2] N. Yoshitani, “Flight trajectory control based on required acceleration for fixed-wing aircraft,” in *Proc. 27th Int. Congr. Aeronautical Sci.*, 2010, vol. 10, pp. 1–10.
- [3] L. Qian, S. Graham, and H. H.-T. Liu, “Guidance and control law design for a slung payload in autonomous landing a drone delivery case study,” *IEEE/ASME Trans. Mechatronics*, vol. 25, no. 4, pp. 1773–1782, Aug. 2020.
- [4] F. Gavilan, R. Vazquez, and E. F. Camacho, “An iterative model predictive control algorithm for UAV guidance,” *IEEE Trans. Aerosp. Elect. Syst.*, vol. 51, no. 3, pp. 2406–2419, Jul. 2015.
- [5] S. Kim, H. Oh, and A. Tsourdos, “Nonlinear model predictive coordinated standoff tracking of a moving ground vehicle,” *AIAA J. Guid. Control Dyn.*, vol. 36, no. 2, pp. 557–566, 2013.
- [6] J. Yang, C. Liu, M. Coombes, Y. Yan, and W.-H. Chen, “Optimal path following for small fixed-wing UAVs under wind disturbances,” *IEEE Trans. Control Syst. Tech.*, vol. 29, no. 3, pp. 996–1008, May 2021.
- [7] D. R. Nelson, D. B. Barber, T. W. McLain, and R. W. Beard, “Vector field path following for small unmanned air vehicles,” in *Proc. IEEE Amer. Control Conf.*, 2006, pp. 5788–5794.
- [8] D. Cabecinhas, C. Silvestre, P. Rosa, and R. Cunha, “Path-following control for coordinated turn aircraft maneuvers,” in *Proc. AIAA Guid., Navigat. Control Conf. Exhibit.*, 2007, pp. 1–19.
- [9] T. Yamasaki, S. Balakrishnan, and H. Takano, “Separate-channel integrated guidance and autopilot for automatic path-following,” *AIAA J. Guid. Control Dyn.*, vol. 36, no. 1, pp. 25–34, 2013.
- [10] Y. Wang, W. Zhou, J. Luo, H. Yan, H. Pu, and Y. Peng, “Reliable intelligent path following control for a robotic airship against sensor faults,” *IEEE/ASME Trans. Mechatronics*, vol. 24, no. 6, pp. 2572–2581, Dec. 2019.
- [11] S. Park, “Design of three-dimensional path following guidance logic,” *Int. J. Aerosp. Eng.*, vol. 2018, 2018, Art. no. 9235124.
- [12] T. Yamasaki, H. Takano, and Y. Baba, “Robust path-following for UAV using pure pursuit guidance,” in *Aerial Vehicles*. London, U.K.: IntechOpen, 2009.

- [13] R. Rysdyk, "Unmanned aerial vehicle path following for target observation in wind," *AIAA J. Guid. Control Dyn.*, vol. 29, no. 5, pp. 1092–1100, 2006.
- [14] N. Cho, Y. Kim, and S. Park, "Three-dimensional nonlinear differential geometric path-following guidance law," *AIAA J. Guid. Control Dyn.*, vol. 38, no. 12, pp. 2366–2385, 2015.
- [15] S. Park, J. Deyst, and J. P. How, "Performance and Lyapunov stability of a nonlinear path following guidance method," *AIAA J. Guid. Control Dyn.*, vol. 30, no. 6, pp. 1718–1728, 2007.
- [16] L. Meier, P. Tanskanen, L. Heng, G. H. Lee, F. Fraundorfer, and M. Pollefeys, "PIXHAWK: A micro aerial vehicle design for autonomous flight using onboard computer vision," *Auton. Robot.*, vol. 33, no. 1/2, pp. 21–39, 2012.
- [17] R. Curry, M. Lizarraga, B. Mairs, and G. H. Elkaim, "L2, an improved line of sight guidance law for UAVs," in *Proc. IEEE Amer. Control Conf.*, 2013, pp. 1–6.
- [18] T. Stastny, "L1 guidance logic extension for small UAVs: Handling high winds and small loiter radii," *CoRR*, vol. abs/1804.0, 2018.
- [19] P. Eng, L. Mejias, X. Liu, and R. Walker, "Automating human thought processes for a UAV forced landing," *J. Intell. Robot. Syst.*, vol. 57, no. 1–4, pp. 329–349, 2010.
- [20] P. Eng, "Path planning, guidance and control for a UAV forced landing," Ph.D. dissertation, School of Engineering Systems, Queensland Univ. Technol., Brisbane, QLD, Australia, 2011.